

AI Job Navigator — Project Schedule and Requirements

Subject document:

AI_Job_Navigator_Client_Requirements_Engineering_Specification_EN_Updated.docx **Date:**
31 August 2026 Version: v5.0

0. Purpose of This Document

This document sets out four things:

1. What we require from your side, and at which point in the project each item is needed
2. The expected cost of development and of running the service
3. The proposed development schedule for the **required** features
4. A separate list of **optional** features, to be scheduled only after the required work is complete

Required and optional

The features described in this document are divided into two groups.


Required features are those defined across the MVP phases of the requirements specification. These form the product, and are the only features in the schedule in Section 3.

Optional features are additions that may be built once every required feature is complete. They are listed in Section 4 with their own estimates, and are deliberately excluded from the Section 3 schedule so that the timeline reflects only what the product needs in order to function.

Development resourcing

Development is carried out by one developer. The timeline in Section 3 is calculated on that basis.

System overview



Job data acquisition policy

We have received your guidance that when acquiring data by scraping, care must be taken not to place a load on the other party's servers. This is reflected in the design.

Collection runs on a schedule with a fixed minimum interval between requests, and is never triggered by user activity. The volume of access to each source is therefore determined by our schedule and stays constant, rather than rising with the number of users. Where a source offers a structured data endpoint we will use it in preference to parsing pages, as it is lighter on their servers and more stable in operation.

Current implementation status

Layer	Status
Conversation infrastructure — response streaming, concurrent access safety, data persistence, conversation history summarization	Complete. Production quality
AI agent framework — function invocation loop, context assembly, AI service abstraction	Complete. Currently connected to one trivial function
Web interface — chat screen, conversation list, theming	Complete for chat. No other screens exist
Account management — registration, login, password reset	Not started
Job data — collection, storage, search, duplicate removal	Not started
Document creation — templates, generation, file output, upload	Not started
Payment, human support, admin dashboard	Not started

Approximately 15 to 20 percent of the total build is complete. What exists is the foundation layer that allows a reliable AI conversation to take place. The features that make the product a career service have not yet been built, and the schedule in Section 3 covers them in full.

1. Required from the Client

1.1 AI service — the one item required immediately

This is the only request in this document that is needed now rather than at a later phase.

The application requires a commercial AI service. We propose **GPT-5.6 Luna**, charged at **USD 0.20 per million input tokens** and **USD 1.20 per million output tokens**.

Why the current development environment is not sufficient. Development is presently running against two stopgap options, and neither can carry the project forward:

Current option	Problem
A locally hosted model	Too slow for practical iteration, and the response quality is not sufficient to judge whether a feature works
A free-tier hosted model	Intermittently blocked by usage limits. Automated testing cannot be relied upon to complete

Which parts of the project depend on it:

Phase	Feature requiring the AI service	Without it
Phase 0	Response quality evaluation framework	No quality baseline can be established
Phase 1	End-to-end technical validation	The architecture cannot be proven to work
Phase 2	AI interview, document content generation	This is the substance of the phase. Cannot be developed or verified
Phase 3	Natural-language condition structuring	Cannot be developed or verified
Phase 4	Japanese response quality work	This is the purpose of the phase
Phase 5	Judgement of when to hand over to a human operator	Cannot be verified

Every phase from Phase 1 onward depends on it.


Cost per operation

Measured against the current system configuration:

Operation	Cost
First message of a conversation	\$0.0010
A normal conversational turn	\$0.0016
A turn that performs a job search	\$0.0033
One conversation history summarization	\$0.0009
A complete 30-turn consultation	approx. \$0.08

A job search costs roughly twice a normal turn because it requires two exchanges with the AI service.

Figure 1 — Example flow: a user asking for job listings



Monthly cost during development

Cost is not constant. It concentrates in the phases where AI behaviour is built and tuned:

Phase	Duration	AI service per month
Phase 0 — Preparation and investigation	1 week	\$8
Phase 1 — Foundation	4 weeks	\$15
Phase 2 — Document creation	4 weeks	\$50
Phase 3 — Job search	5 weeks	\$45
Phase 4 — Quality and refinement	3 weeks	\$40
Phase 5 — Commercial functions	5 weeks	\$10
Phase 6 — Hardening and launch	3 weeks	\$20
Total across the project	25 weeks	approx. USD 180

We propose a ceiling of USD 50 per month, expected to be reached only during Phases 2 to 4.

For context: reaching USD 100 in a month would require approximately 4,400 full-length AI requests every day. Manual testing by one developer cannot approach that volume, so the limit functions as a safety measure rather than an expected cost.

The principal risk is a software defect rather than usage. A fault causing the AI to call the same function repeatedly would consume approximately USD 5 per hour if left unattended. Two safeguards will be implemented in Phase 0: a hard spending limit on the service account, and a maximum-iteration guard within the agent itself.

1.2 Items required at later phases

None of the following is needed now. Each is listed against the phase at which it becomes necessary, so that it can be prepared in advance rather than requested at short notice.

Required by	Item	Purpose
Phase 2	Five to ten real 履歴書 / 職務経歴書, with personal details removed	Three separate purposes — see below
Phase 2	The 履歴書 / 職務経歴書 format to be used as the base	Document template development
Phase 2 onward	A nominated Japanese-speaking reviewer	Ongoing review of AI-generated Japanese. See 1.3
Phase 5	Payment system credentials and integration details	See 1.4
Phase 6	Production domain and DNS control	TLS certificates and the production environment

On the request for real documents. These are needed for three distinct reasons, and no invented example can substitute for any of them:

1. **To build the template.** We need to see the actual layout of the format in use — which sections appear, in what order, under what headings, with what spacing. A fill-in template cannot be built from imagination.
2. **To show the AI what good output looks like.** Providing the AI with real, approved examples is the most effective way to make its Japanese read naturally. This is the direct remedy for the register issue described in Section 1.3.
3. **To measure quality.** Judging whether the AI's output is acceptable requires something to compare it against. Without real examples, the question of whether output is good enough has no answer.

Real documents also carry variety that invented ones do not — career gaps, career changes, part-time periods — which is exactly what the system must handle correctly.

Personal details are not needed and should be removed. What is required is the structure and the writing style.

1.3 AI response quality review — begins at Phase 2

This is a request for ongoing involvement rather than a one-off deliverable, which is why it is raised early even though it is not needed until Phase 2.

The AI model produces grammatically correct Japanese. The risk is not correctness but **register** — whether a 職務経歴書 reads as something a Japanese hiring manager or company officer would find appropriate. This cannot be measured by any automated benchmark.

We therefore ask that a Japanese-speaking reviewer be nominated on your side from Phase 2 onward, to review sample AI output at each demonstration and indicate which responses are preferable. Approximately two hours per demonstration cycle is expected.

How that feedback is used. Reviewer feedback is applied in escalating steps: first by adjusting the AI's written instructions, then by adding approved examples for the AI to follow, and if necessary by retrieving relevant approved examples automatically for each case. Only if those measures reach their limit and quality is still short of acceptable would we recommend moving to a higher-grade AI model. That is a configuration change requiring no modification to the application, though the running cost rises by approximately a factor of ten to fifteen — from approximately USD 0.08 per consultation to approximately USD 0.80–1.20 per consultation. At ten users this remains modest; we would recommend reassessing before any significant increase in user numbers.

1.4 Payment system — required at Phase 5

We understand the payment system to be the environment indicated at the reference provided, which is a GMO Payment Gateway environment rather than a general-purpose international provider.

To size and build this work we will need the following, by Phase 5:

Item	Note
Which integration method is to be used	GMO offers a hosted redirect type (リンクタイプ), a server module type (モジュールタイプ), and a direct protocol type (プロトコルタイプ). The effort involved differs substantially between them
Test environment credentials (shop ID, shop password)	Issued under your GMO contract. No payment code can be written or tested before these are available
Which payment methods are in scope	Credit card only, or also convenience store payment, bank transfer, carrier billing. Each is a separate integration, and non-card methods complete asynchronously rather than immediately
Whether 3-D Secure is required	Adds an additional bank verification

Item	Note
	step to the payment flow
Pricing plans, free-tier scope, and cancellation rules	Still listed as undecided in the specification

The Phase 5 estimate is provisional until the integration method is known.

A hosted redirect integration is approximately two weeks; a direct protocol integration is approximately four. The schedule in Section 3 assumes an intermediate figure.

2. Costs

All figures are in US dollars and assume the AWS Asia Pacific (Tokyo) region. They are estimates and should be confirmed with the AWS Pricing Calculator before being treated as final.

2.1 Development approach — local first

Development is carried out on a local machine, and cloud infrastructure is provisioned only when a feature requires it. Provisioning a full cloud environment from the first week would mean paying for it months before it is needed.

The points at which cloud infrastructure becomes necessary:

Requirement	Phase	Reason
Unattended scheduled collection	Phase 1	A development machine is not continuously available; the job corpus must accumulate without interruption
Payment gateway callbacks	Phase 5	The payment provider must be able to reach the application server directly. A local machine cannot receive these
TLS, production domain, session security	Phase 6	Required before any public use
Load testing at ten concurrent users	Phase 6	Cannot be measured locally with confidence

A single deployment is also carried out within Phase 1 to confirm that the application deploys correctly, so that any deployment issue is identified at the start of the project rather than at the end.

2.2 Development period cost

Phase	Duration	Environment	AWS	AI service	Monthly rate
Phase 0 — Preparation and investigation	1 week	Local only	\$0	\$8	\$8
Phase 1 — Foundation	4 weeks	Local + collection instance	\$8	\$15	\$23
Phase 2 — Document creation	4 weeks	Local + collection instance	\$8	\$50	\$58
Phase 3 — Job search	5 weeks	Local + collection instance	\$10	\$45	\$55
Phase 4 — Quality and refinement	3 weeks	Local + collection instance	\$10	\$40	\$50
Phase 5 — Commercial functions	5 weeks	Staging provisioned	\$40	\$10	\$50
Phase 6 — Hardening and launch	3 weeks	Production + staging	\$85	\$20	\$105

Total cost of development across the entire project: approximately USD 330.

Of this, approximately USD 150 is AWS and USD 180 is the AI service.

Development servers are assumed to run during working hours only and to be stopped overnight and at weekends, which reduces their cost by approximately seventy percent against continuous operation. If they were instead left running continuously, the AWS portion would rise to approximately USD 450 for the project, while the AI portion would be largely unchanged, since AI cost is determined by the volume of testing performed rather than by elapsed time.


2.3 Production at launch

Assumes ten concurrent users, as specified.

Item	Specification	Monthly
Application server	EC2 t4g.small (2 vCPU / 2 GB)	\$18

Item	Specification	Monthly
Database	RDS PostgreSQL db.t4g.small, 50 GB, single-AZ	\$45
Job data collection	Fargate scheduled task, approx. 24 hrs/month	\$3
Document storage	S3	\$1
Frontend delivery	CloudFront + S3	\$1
DNS	Route 53	\$0.50
Logging	CloudWatch	\$3
Data transfer		\$2
Infrastructure subtotal		\$74
AI service	10 concurrent users	\$5
Total		approx. USD 79 / month

Production monthly cost, USD



Adding a load balancer for redundancy raises this by approximately USD 20 per month. It is not recommended at ten concurrent users, and is the natural first upgrade once traffic grows.

2.4 On reserved pricing

AWS offers discounts of approximately 30 to 40 percent in exchange for a one-year commitment to a given level of usage.

We recommend against committing at this stage. At ten concurrent users the total infrastructure cost is approximately USD 74 per month, so the saving would be around USD 25. Securing it would require fixing the server configuration for a year, before any real usage has been observed, and while user numbers are expected to change.

We propose running on standard on-demand pricing for the first three to six months after launch, measuring actual usage, and purchasing a one-year commitment against the observed baseline at that point.

2.5 Two cost items to be aware of

NAT Gateway. A commonly used AWS network component charged at approximately USD 32 per month regardless of usage. Many standard reference architectures include it by default. At this scale it is unnecessary, and the design proposed here avoids it. Were it included, it would be the largest single line item in Section 2.3.

Database storage growth. Server capacity will not be strained by ten users, as the AI service performs the processing. Storage will grow: approximately 100,000 job records with search indexes require around 600 MB for the index alone. We recommend budgeting for 50 to 100 GB within the first year. Retaining uploaded and generated documents indefinitely, as currently intended, adds to this over time and is worth reviewing annually.

3. Development Schedule — Required Features

3.1 Method

Development proceeds in seven phases. **Each phase concludes with a demonstration of working software and a written progress report.**

Phase length varies according to the work involved rather than being fixed, since the phases differ considerably in size. Automated testing is carried out within each phase rather than deferred to a single stage at the end.

3.2 Scope of the initial release

The following are delivered in a defined initial form at launch, with the remainder addressed afterwards. This is set out explicitly so that the scope behind the timeline is clear, and so that any item can be brought forward if you would prefer it included.

Item	At launch	Addressed afterwards
Job sources	Two sources	Remaining two sources added
Duplicate removal	Exact matching on company, title and location	Cross-source matching, once a large body of real data exists
Document format	One 履歷書 / 職務經歷	Additional formats and Word output

Item	At launch	Addressed afterwards
	書 format, PDF output	
Admin dashboard	The functions required for operation	Remaining functions, as operational need appears

The principle applied is to limit breadth rather than depth. Additional sources and additional templates are straightforward to add to a working system. Duplicate removal is the one item that becomes harder to address the longer it is deferred, which is why it is reduced rather than omitted, and why data collection begins in the first working phase.

3.3 Two principles in the plan

Job data collection begins in Phase 1 and runs continuously thereafter.

Duplicate removal depends on having a substantial body of real data, and that data accumulates over calendar time rather than through effort. By the time duplicate removal is built in Phase 3 and refined in Phase 4, approximately two to three months of collected data will be available to work against. Raw pages are stored before any field extraction, so that when extraction logic is improved it can be re-run over data already held rather than requiring the sources to be accessed again.

Phase 1 concludes with an end-to-end technical validation. Before the individual features are built out, a complete path through the system is proven to work: a user logs in, asks a question, the AI invokes the search function, and real job data is returned. This is deliberately minimal and is not the finished product. Its purpose is to confirm that the architecture holds together, and to identify any integration problem at the beginning of the project rather than at the end, when remedy is most costly.

3.4 The job data pipeline

Figure 2 — Job data flow, from collection through to presentation



3.5 Schedule

Phase	Content	Duration
Phase 0	Preparation and investigation Local development environment AI service configuration, spending limit, iteration guard Japanese quality evaluation framework Investigation of the	1 week

Phase	Content	Duration
Phase 1	job sources	4 weeks
	Foundation	
	Account management: registration, login, sessions, password reset, email verification	
	Security hardening of the existing interface	
Phase 2	Job data structure and database design	4 weeks
	Collection framework with request rate control	
	First source collecting continuously	
	End-to-end technical validation	
Phase 3	Deployment verification	5 weeks
	Document creation	
	履歷書 / 職務經歷書 template and generation	
	PDF output and download	
Phase 4	AI interview refinement	3 weeks
	Second source added to collection	
	Job search	
	Normalization of salary units, place names, holiday terminology	
	Duplicate removal	
	Natural-language condition structuring	
	Search function and AI integration	
	Document upload and text extraction	
Phase 4	Quality and refinement	3 weeks
	Japanese response quality, with reviewer	

Phase	Content	Duration
	feedback	
	Posting currency verification, acquisition date display	
	User confirmation screen, result presentation	
	Duplicate removal tuning	
Phase 5	Commercial functions	5 weeks
	Payment integration, plans, cancellation	
	In-app human support and handover	
	Admin dashboard	
Phase 6	Hardening and launch	3 weeks
	Performance, security, load testing at 10 concurrent users	
	Production environment, acceptance testing support	
Total		25 weeks
Approximately six months.		





Figure 3 — Phase sequence. Each phase concludes with a demonstration.



3.6 On the original three-month estimate

The original three-month figure was produced before the requirements specification existed, and covered a smaller scope. The specification as written represents approximately twice that work.

The increase comes from work that was not present in the original outline: account management, payment processing, the administrative dashboard, human support, security hardening, acceptance testing, and the removal of duplicate job listings appearing across multiple sources.

3.7 Schedule confirmation

Two items cannot be estimated with full confidence until real data has been examined: the stability of collection from the job sources, and the accuracy achievable in duplicate removal. A third, the payment integration, depends on the method described in Section 1.4.

We therefore ask that the schedule above be treated as provisional, and that firm dates be agreed at the end of Phase 0. Phase 0 includes investigation of the job sources, and at its conclusion — one week from project start — a confirmed schedule can be provided.

Where a phase runs longer than estimated, this will be reported at the phase boundary rather than at the end of the project, and the scope items in Section 3.2 remain available as adjustments in either direction.

If a shorter timeline is required, the most effective single change is to defer Phase 5 — payment, human support, and the admin dashboard — which brings a complete and demonstrable product forward by five weeks.

4. Optional Features

The following are **not included in the schedule in Section 3** and are not required for the product to function. They are recorded here with estimates so that they can be scheduled deliberately, as a separate piece of work, once every required feature is complete.

Feature	Duration	Depends on
Notification system — saved search conditions, daily matching against new	3 weeks	Phase 3

Feature	Duration	Depends on
postings, email delivery infrastructure, notification preferences		
Remaining two job sources	2 weeks	Phase 1 collection framework
Refined duplicate removal — cross-source matching rather than exact matching	3 weeks	A large body of real data
Career Advancement Assessment — assessment of a completed 職務經歷書, with administrator review	3 weeks	Phase 2 and the admin dashboard
LINE integration — official account, Messaging API, and linking a LINE account to a service account	2.5 weeks	Notification system
OCR — reading scanned images and photographs of documents	2 weeks	Phase 3
Full admin dashboard — the remaining functions	2 weeks	Phase 5
Additional document formats and Word output	1 week	Phase 2
Document delivery by email	3 days	Notification system
Social media job sources	3 weeks	Investigation first; may prove infeasible
Total, if all are built	approx. 22 weeks	

These are listed in the order in which we would recommend building them. None needs to be decided now.

4.1 Notes on individual items

The notification system is the item most likely to be wanted first. It is the reason saved search conditions appear in the specification. Until it is built, a user re-runs a search themselves rather than being informed of new matches. The product functions without it; it is less convenient.

Email is not otherwise required. With notifications excluded from the required scope, generated documents are delivered by download from the browser. A minimal email capability is still built in Phase 1 for account functions — password reset and address verification — as an account system without them is not usable in practice. That is infrastructure supporting accounts, and is distinct from the notification feature.

Social media sources may prove infeasible rather than merely costly. For several platforms, no means of searching public posts is made available to third parties. The estimate above assumes a route exists, and investigation should precede any commitment.

Career Advancement Assessment cannot be built earlier regardless of priority. It operates on a completed 職務経歴書 and requires an administrator review flow, so it depends on both document creation and the admin dashboard.

5. Assumptions

Where any of these proves incorrect, the affected figures will change.

- One developer for the duration of the project
 - Ten concurrent users at launch
 - Two job sources at launch, with the remaining two added afterwards
 - Duplicate removal by exact matching at launch, refined once sufficient real data exists
 - One 履歴書 / 職務経歴書 format with PDF output at launch; additional formats and Word output added later
 - Admin dashboard limited at launch to the functions required for operation
 - Everything listed in Section 4 is excluded from the schedule in Section 3
 - Human support via in-app chat only
 - Generated documents delivered by download; email is used only for account functions
 - Uploaded and generated documents retained indefinitely, per current intent
 - Development conducted locally, with cloud infrastructure provisioned per Section 2.1
 - Phase length varies according to the work involved
 - User interface and screen design determined by us, subject to your review
 - The Terms of Service and Privacy Policy are produced by your side
-

6. Next Steps

Confirmation of this
document



Approval of the AI service
budget
Section 1.1



Phase 0 begins
one week



.. ..

The only item required in order to begin is the AI service approval in Section 1.1. Everything else in Section 1 is needed at a later phase and is listed so that it can be prepared in good time.

We look forward to your confirmation.